

DEEP LEARNING FOR SPEECH & LANGUAGE

Winter Seminar UPC TelecomBCN, 24 - 31 January 2017



Instructors



Antonio Bonafonte J. Adrián Rodríguez Fonollosa Marta R. Costa-jussà Javier Hernando Santiago Pascual Elisa Sayrol Xavier Giró

Organizers



Image Processing Group
Signal Theory and Communications Department



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

+ info: [TelecomBCN.DeepLearning.Barcelona](https://www.telecombcn.com/deeplearning-barcelona)

[\[course site\]](#)

Day 1 Lecture 2

The Perceptron



Santiago Pascual



Outline

1. Supervised learning: Regression/Classification
2. Linear regression
3. Logistic regression
4. The Perceptron
5. Multi-class classification
6. The Neural Network
7. Metrics

Machine Learning techniques

We can categorize three types of learning procedures:

1. **Supervised Learning:**

$$y = f(x)$$



2. **Unsupervised Learning:**

$$f(x)$$

3. **Reinforcement Learning:**

$$y = f(x)$$

z

We have a labeled dataset with pairs (\mathbf{x}, \mathbf{y}) , e.g. classify a signal window as containing speech or not:

$\mathbf{x1} = [x(1), x(2), \dots, x(T)]$

$\mathbf{y1} = \text{"no"}$

$\mathbf{x2} = [x(T+1), \dots, x(2T)]$

$\mathbf{y2} = \text{"yes"}$

$\mathbf{x3} = [x(2T+1), \dots, x(3T)]$

$\mathbf{y3} = \text{"yes"}$

...

Supervised Learning

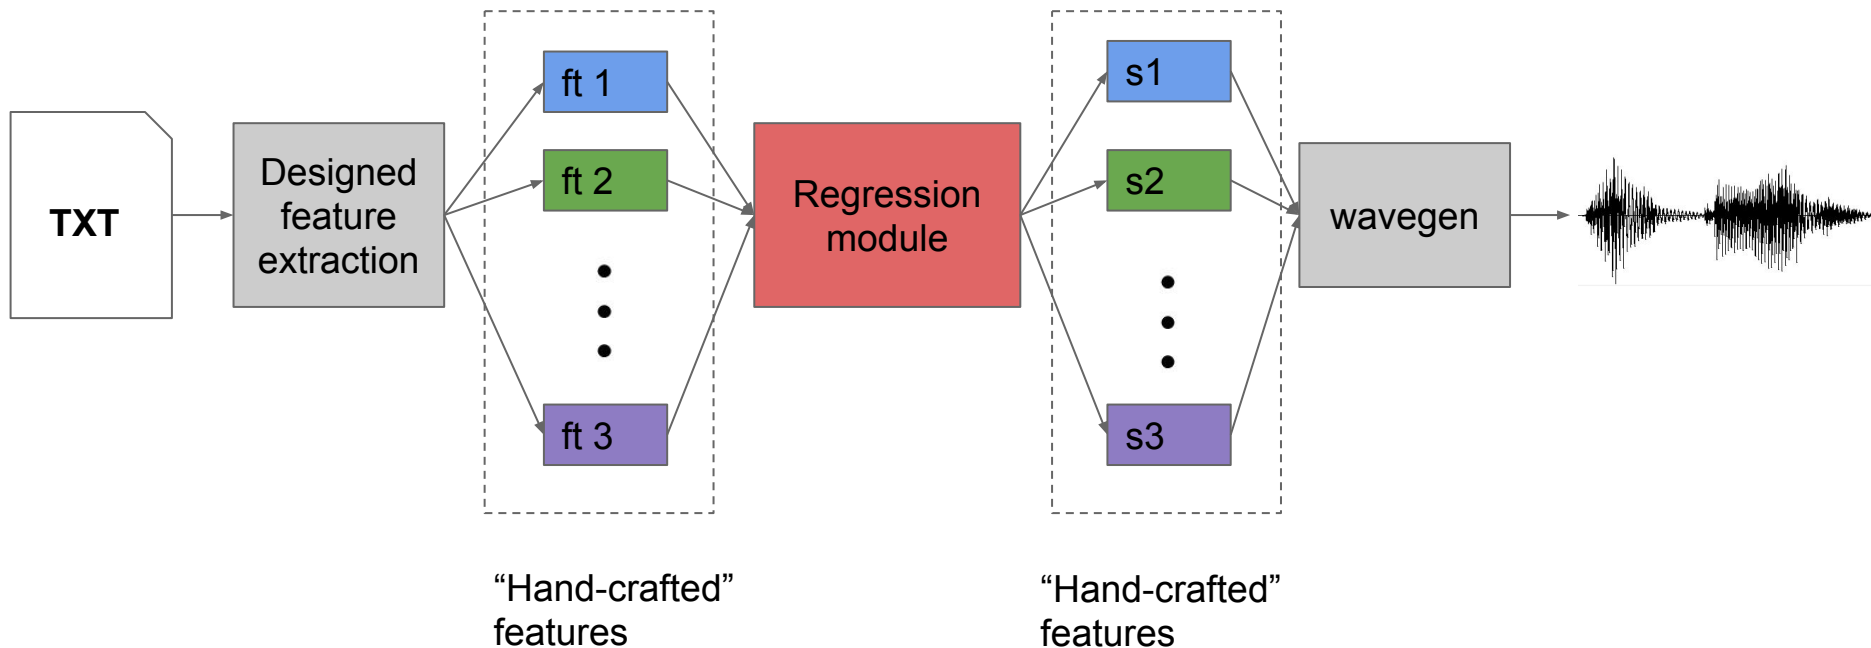
Build a function: $y = f(\mathbf{x})$, $\mathbf{x} \in \mathbb{R}^m$, $y \in \mathbb{R}^n$

Depending on the type of outcome we get...

- Regression: y is continuous (e.g. temperature samples $y = \{19^\circ, 23^\circ, 22^\circ\}$)
- Classification: y is discrete (e.g. $y = \{1, 2, 5, 2, 2\}$).
 - Beware! These are unordered categories, not numerically meaningful outputs: e.g. `code[1] = "dog"`, `code[2] = "cat"`, `code[5] = "ostrich"`, ...

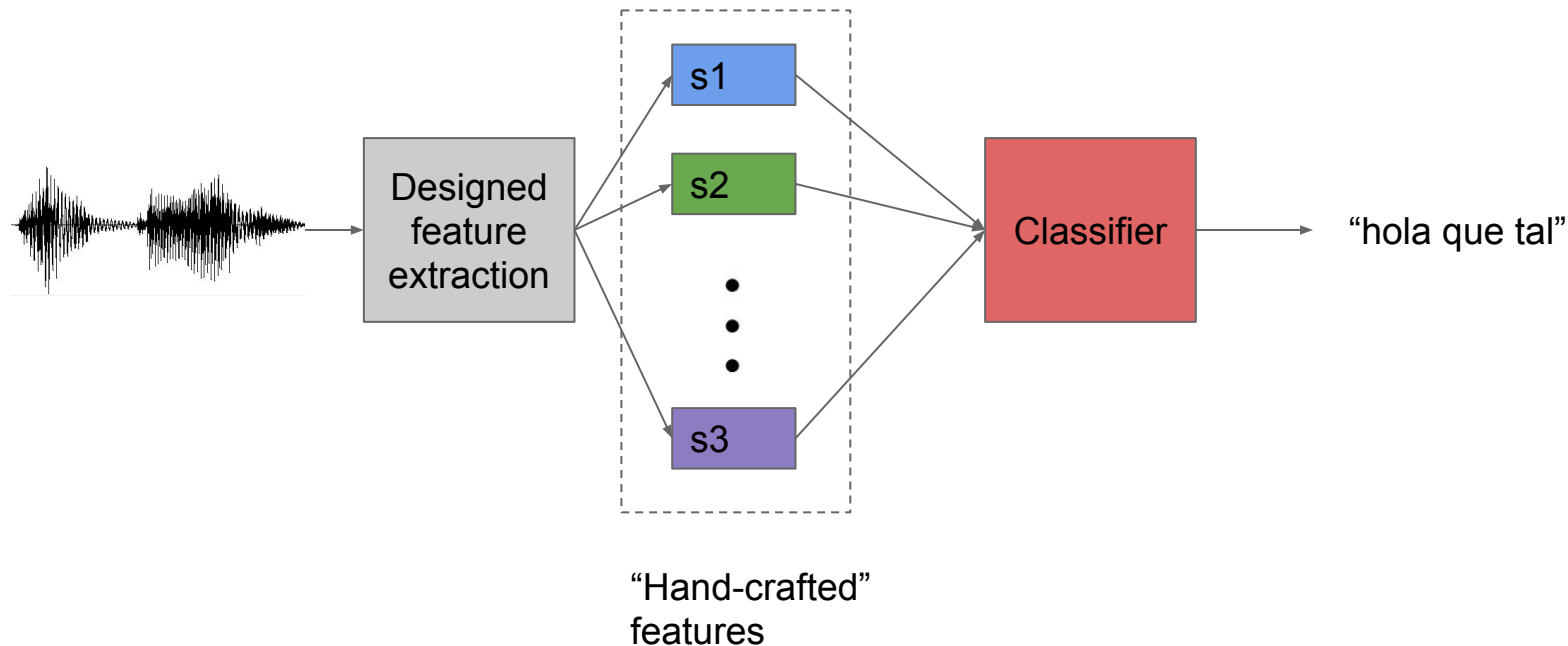
Regression motivation

Text to Speech: Textual features \rightarrow Spectrum of speech (many coefficients)



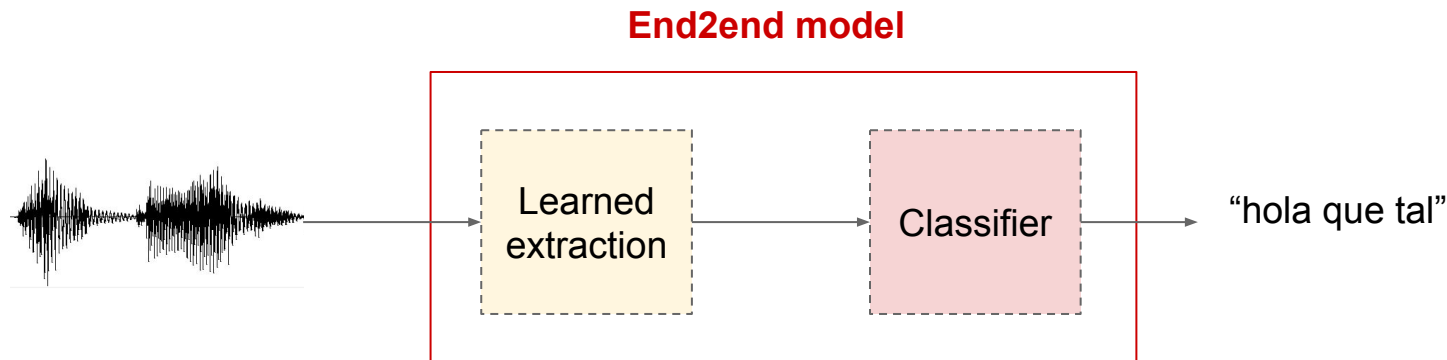
Classification motivation

Automatic Speech Recognition: Acoustic features → Textual transcription (words)



What “deep-models” means nowadays

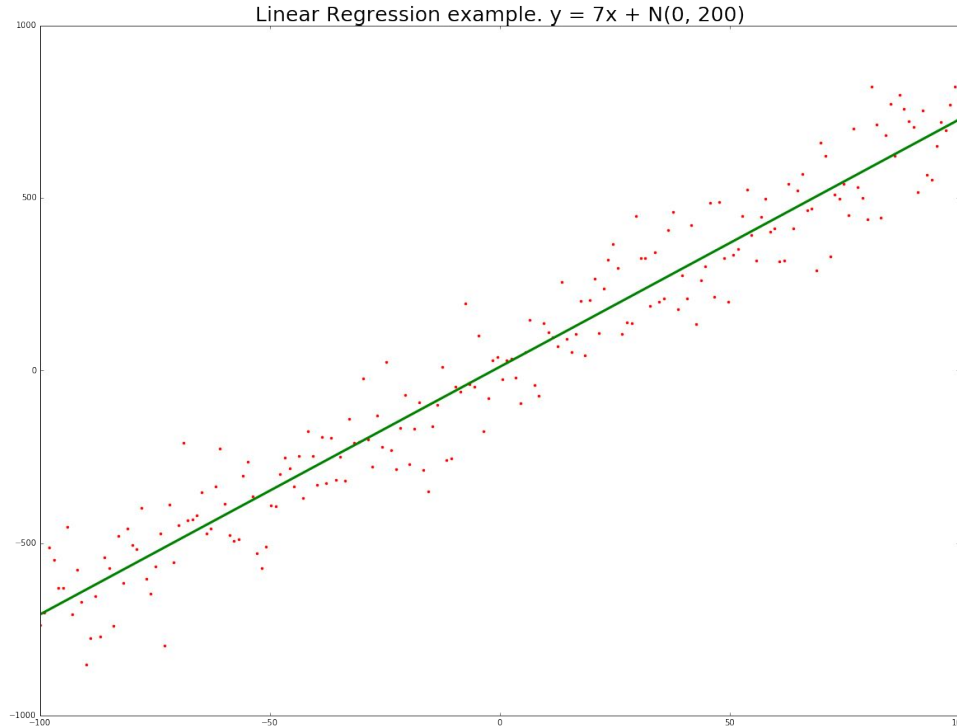
Learn the representations as well, not only the final mapping → **end2end**



Model maps raw inputs to raw outputs, no intermediate blocks.

Linear Regression

Function approximation $y = \omega \cdot x + \beta$, with learnable parameters $\theta = \{\omega, \beta\}$



Linear Regression

We can also make the function more complex for \mathbf{x} being an M-dimensional set of features: $y = \omega_1 \cdot x_1 + \omega_2 \cdot x_2 + \omega_3 \cdot x_3 + \dots + \omega_M \cdot x_M + \beta$

e.g. we want to predict the price of a house based on:

x_1 = square-meters (sqm)

$x_{2,3}$ = location (lat, lon)

x_4 = year of construction (yoc)

price = $\omega_1 \cdot (\text{sqm}) + \omega_2 \cdot (\text{lat}) + \omega_3 \cdot (\text{lon}) + \omega_4 \cdot (\text{yoc}) + \beta$

- Fitting $f(\mathbf{x})$ means adjusting (**learning**) the values $\theta = \{\omega_1, \omega_2, \dots, \omega_M, \beta\}$
 - How? Will see in training chapter, stay tuned!

Logistic Regression

In the classification world we talk about **Probabilities**, and more concretely:

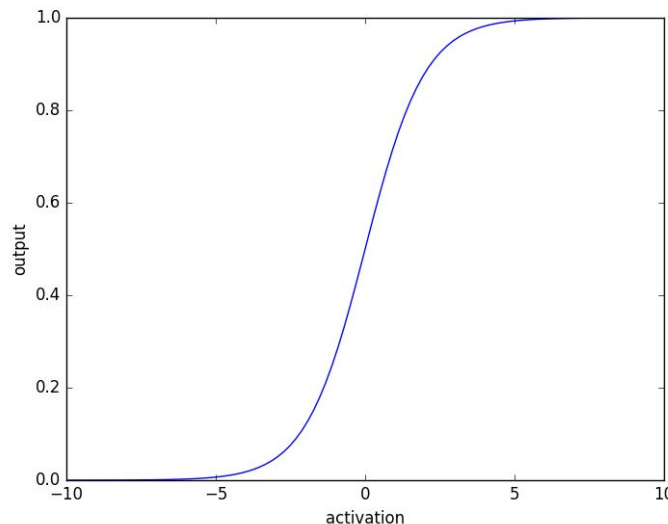
Given \mathbf{x} input data features \rightarrow Probability of y being:

- a dog $P(y=\text{dog}|\mathbf{x})$
- a cat $P(y=\text{cat}|\mathbf{x})$
- a horse $P(y=\text{horse}|\mathbf{x})$
- whatever $P(y=\text{whatever}|\mathbf{x})$.

We achieve so with the **sigmoid function**!

$$\sigma(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{w}^T \cdot \mathbf{x} + b}}$$

Note: This is a binary classification approach

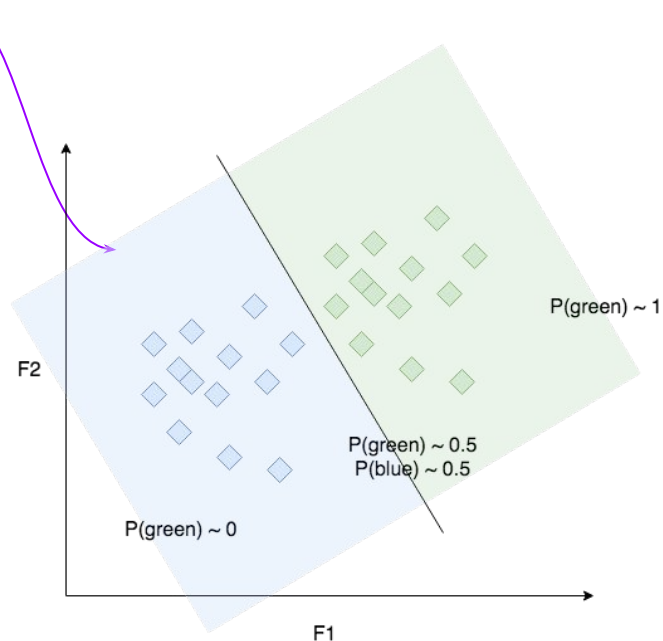
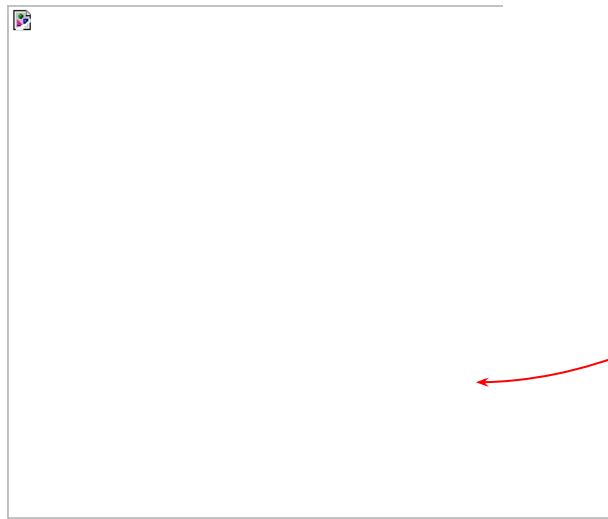


Bounded $\sigma(\mathbf{x}) \in (0, 1)$

Logistic Regression

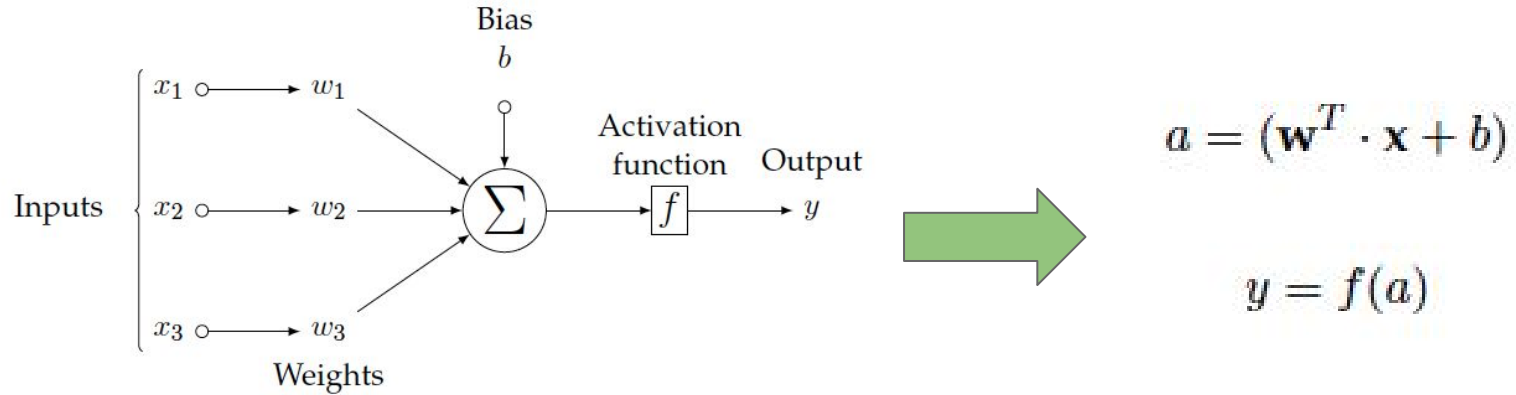
Interpretation: build a delimiting boundary between our data classes + apply the sigmoid function to estimate a probability in every point in the space.

$$\sigma(x) = \frac{1}{1 + e^{-w^T \cdot x + b}}$$



The Perceptron

Both operations, linear regression and logistic regression, follow the scheme in the Figure:



Depending on the Activation function f we have a linear/non-linear behavior:

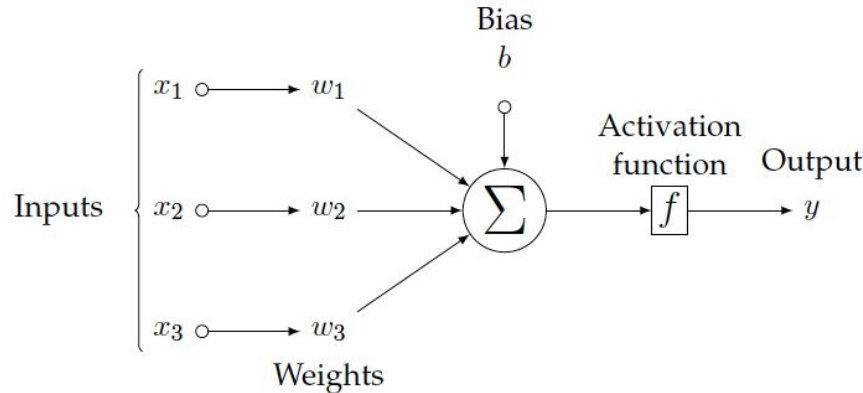
if $f == \text{identity} \rightarrow \text{linear regression}$

if $f == \text{sigmoid} \rightarrow \text{logistic regression}$

The Perceptron

The output is then derived by a **weighted sum** of the inputs **plus a bias term**.

Weights and bias are the parameters we keep (once learned) to define a neuron.



$$a = (\mathbf{w}^T \cdot \mathbf{x} + b)$$

$$y = f(a)$$

The Perceptron

Actually the artificial neuron is seen as an analogy to a biological one.

Real neuron fires an impulse once the sum of all inputs is over a threshold.

The sigmoid emulates the thresholding behavior → act like a switch.

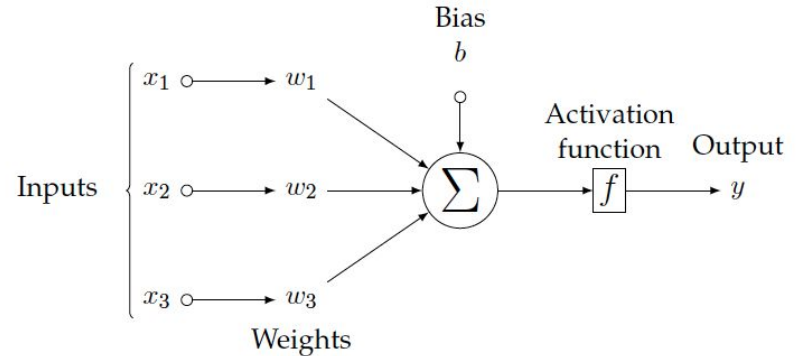
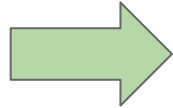
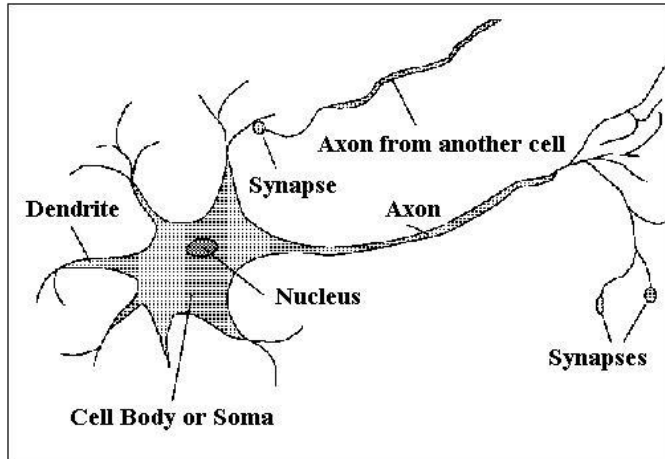
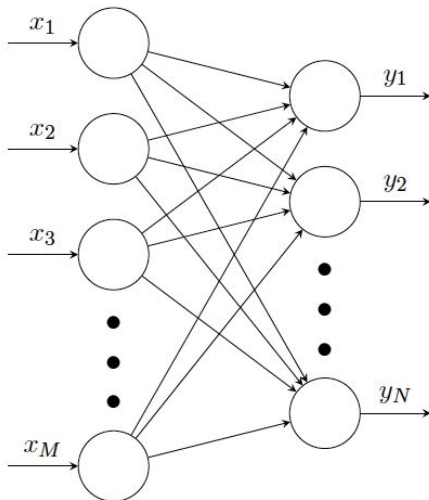
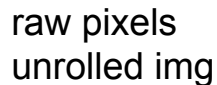


Figure credit: [Introduction to AI](#)

Multi-class classification

Natural extension: put many neurons in parallel, each processing its binary output out of N possible classes.



0.3 “dog”

Page 10

0.08 “cat”



●

•

●

Page 10 of 10

0.6 “whatever”

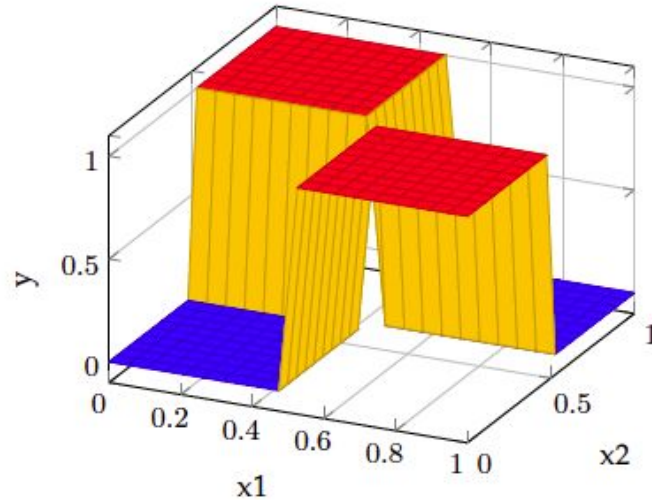
Softmax function

$$P(y = k|\mathbf{x}) = \frac{\exp \mathbf{x}^T \mathbf{w}_k}{\sum_{n=1}^N \exp \mathbf{x}^T \mathbf{w}_n}$$

Normalization factor,
remember: we want a pdf at
the output! \rightarrow all output P's
sum up to 1.

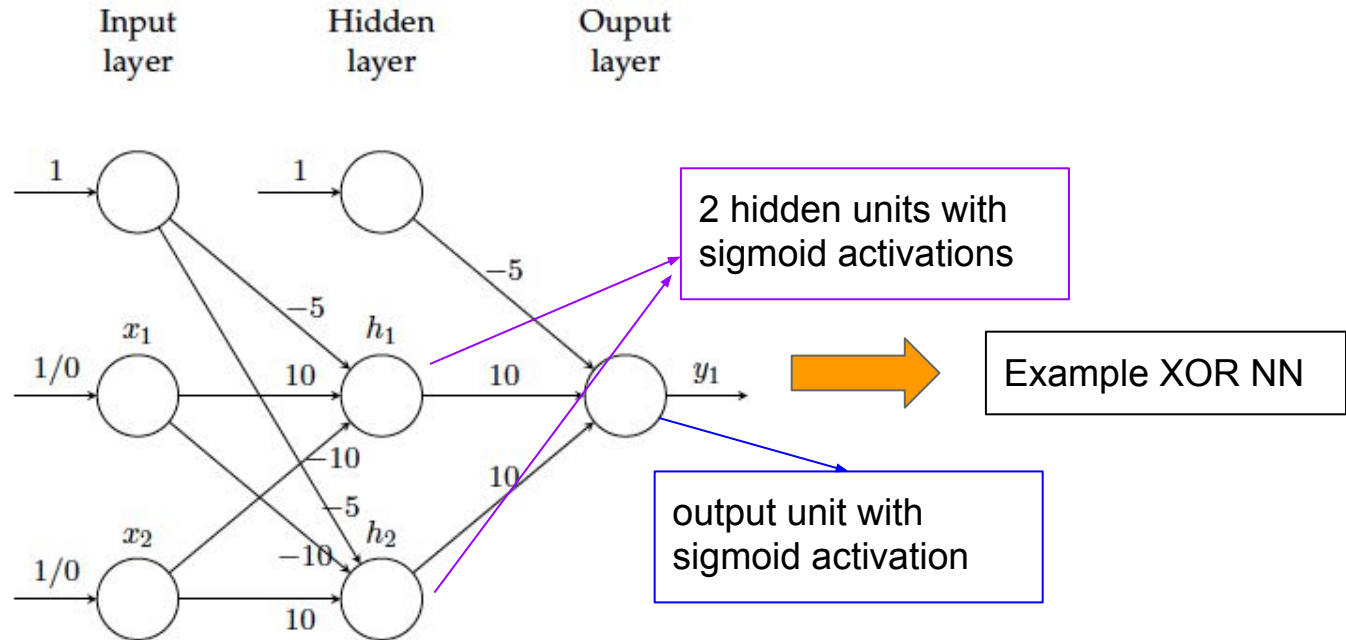
The Neural Network

The XOR problem: sometimes a single neuron is not enough → Just a single decision split doesn't work



The Neural Network

Solution: arrange many neurons in a first intermediate **non-linear** mapping (Hidden Layer), **connecting everything** from layer to layer in a **feed-forward** fashion.



Warning! Inputs are not neurons, but they are usually depicted like neurons.

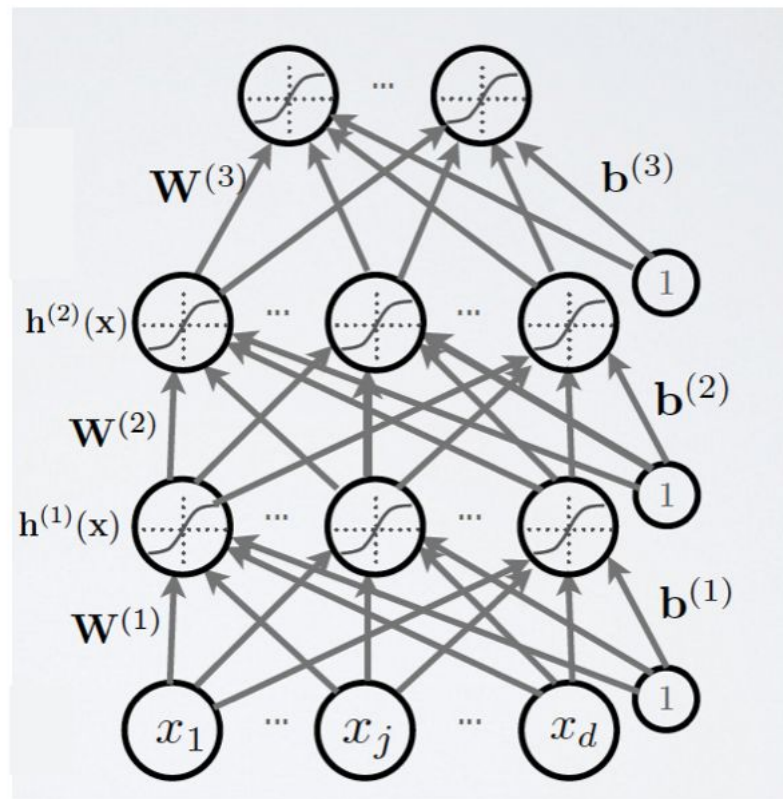
The Neural Network

The i -th layer is defined by a matrix \mathbf{W}_i and a vector \mathbf{b}_i , and the activation is simply a dot product plus \mathbf{b}_i :

$$h_i = f(\mathbf{W}_i \cdot h_{i-1} + b_i)$$

Num parameters to learn at i -th layer:

$$N_{params}^i = N_{inputs}^i \times N_{units}^i + N_{units}^i$$



Slide Credit: Hugo Laroché NN course

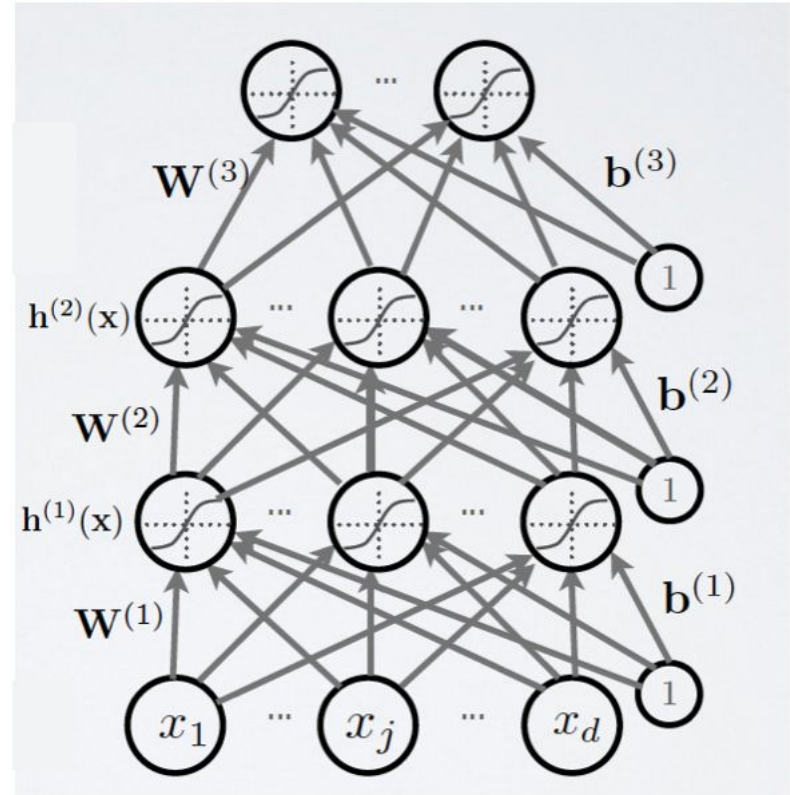
The Neural Network

The i -th layer is defined by a matrix \mathbf{W}_i and a vector \mathbf{b}_i , and the activation is simply a dot product plus \mathbf{b}_i :

$$h_i = f(\mathbf{W}_i \cdot h_{i-1} + \mathbf{b}_i)$$

Num parameters to learn at i -th layer:

$$N_{params}^i = N_{inputs}^i \times N_{units}^i + N_{units}^i$$



Slide Credit: Hugo Laroché NN course

The Neural Network

Important remarks:

- We can put as many hidden layers as we want whenever training can be effectively done and we have enough data (next chapters)
 - **The amount of parameters to estimate grows very quickly** with the num of layers and units! → There is **no formula to know the amount of units per layer nor the amount of layers**, pitty...
- **The power of NNets comes from non-linear mappings**: hidden units must be followed by a non-linear activation!
 - *sigmoid, tanh, relu, leaky-relu, prelu, exp, softplus, ...*

Regression metrics

In regression the metric is chosen based on the task:

- For example in TTS there are different metrics for the different predicted parameters:
 - Mel-Cepstral Distortion, Root Mean Squared Error F0, duration, ...

Classification metrics

Confusion matrices provide a by-class comparison between the results of the automatic classifications with ground truth annotations.

		Automatic		
		class1	class2	class3
Manual	class1	12	1	0
	class2	3	13	0
	class3	0	0	20

		Automatic		
		class1	class2	class3
Manual	class1	100%	0%	0%
	class2	0%	100%	0%
	class3	0%	0%	100%

Classification metrics

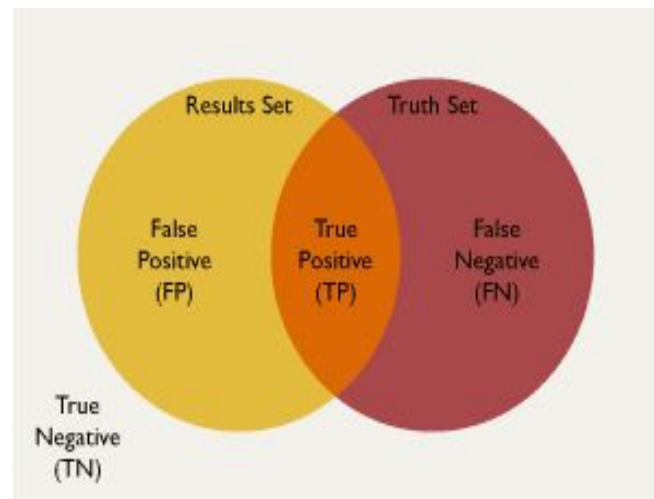
Correct classifications appear in the diagonal, while the rest of cells correspond to errors.

		Prediction		
		Class 1	Class 2	Class 3
Ground Truth	Class 1	x(1,1)	x(1,2)	x(1,3)
	Class 2	x(2,1)	x(2,2)	x(2,3)
	Class 3	x(3,1)	x(3,2)	x(3,3)

Classification metrics

Special case: Binary classifiers in terms of “Positive” vs “Negative”.

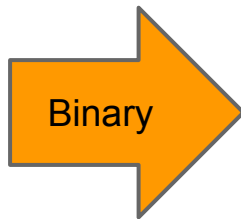
		Prediction	
		Positives	negative
Ground Truth	Positives	True positive (TP)	False negative (FN)
	negative	False positives (FP)	True negative (TN)



Classification metrics

The “accuracy” measures the proportion of correct classifications, not distinguishing between classes.

$$Accuracy = \frac{\sum_{i=1}^3 x(i, i)}{\sum_{i=1}^3 \sum_{j=1}^3 x(i, j)}$$



$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

		Prediction		
		Class 1	Class 2	Class 3
Ground Truth	Class 1	x(1,1)	x(1,2)	x(1,3)
	Class 2	x(2,1)	x(2,2)	x(2,3)
	Class 3	x(3,1)	x(3,2)	x(3,3)

		Prediction	
		Positives	negative
Ground Truth	Positives	True positive (TP)	False negative (FN)
	Negative	False positives (FP)	True negative (TN)

Classification metrics

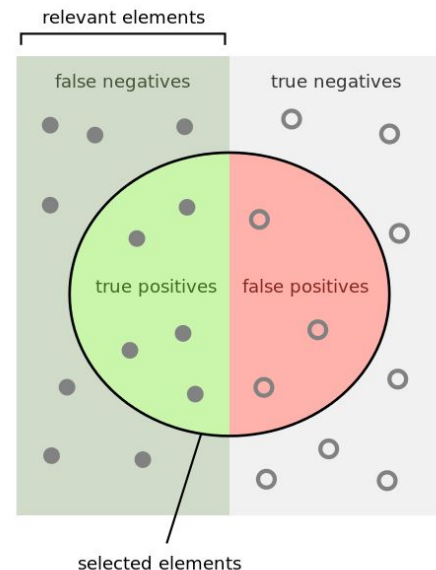
Given a reference class, its Precision (P) and Recall (R) are complementary measures of relevance.

Example: Relevant class is “Positive” in a binary classifier.

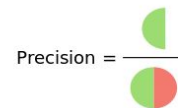
		Prediction	
		Positives	Negatives
Ground Truth	Positives	True positive (TP)	False negative (FN)
	Negatives	False positives (FP)	

$$Recall = \frac{TP}{TP + FN}$$

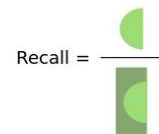
$$Precision = \frac{TP}{TP + FP}$$



How many selected items are relevant?



How many relevant items are selected?



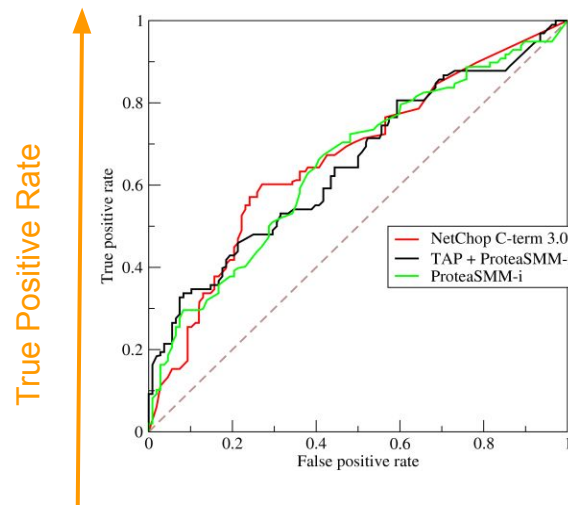
Classification metrics

Binary classification results often depend from a parameter (eg. decision threshold) whose value directly impacts precision and recall.

For this reason, in many cases a Receiver Operating Curve (ROC curve) is provided as a result.

$$\text{True Positive Rate} = \frac{TP}{TP + FN} = \text{Recall} = \text{Sensitivity}$$

$$\text{False Positive Rate} = \frac{FP}{TP + FN} = 1 - \text{specificity}$$



Thanks ! Q&A ?



[@Santty128](https://twitter.com/Santty128)